MYE023: Homework #1

Due on Monday, April 3, 2017

Vassilios V. Dimakopoulos

George Z. Zachos

April 2, 2017

Contents

1 Exercise #1

1.1 About

This exercise is about the calculation of the mathematical constant π using `POSIX` threads and dynamic scheduling. During dynamic scheduling the parallelizable loops are divided into chunks of iterations (tasks) and are dispatched to the threads available to the runtime system for execution. The dispatch takes place in respect to the current processor workload where each thread executes and as a result load balancing is achieved. In case chunk size is one (1) iteration, we refer to this technique as self-scheduling. The purpose of this exercise is to time the calculation of π and observe how altering the number of threads will affect execution time for a given chunk size.

1.2 Experiment details

The calculation consists of $5 * 10^8$ loop iterations, while thread number takes value in $\{1, 4, 16\}$ and chunk size in $\{1, 10, 10^2, 10^3, 10^4, 10^5\}$.

1.2.1 System Specifications

The experiments were conducted on a Dell OptiPlex 7020:

- CPU: Intel® Core™ i5-4590 CPU @ 3.30GHz (64 bit)

- RAM: 2 DIMMs x4GiB @ 1600MHz DDR3

- Cache line size: 64B (in all levels)

- Cache associativity:

    - L1, L2: 8-way set associative
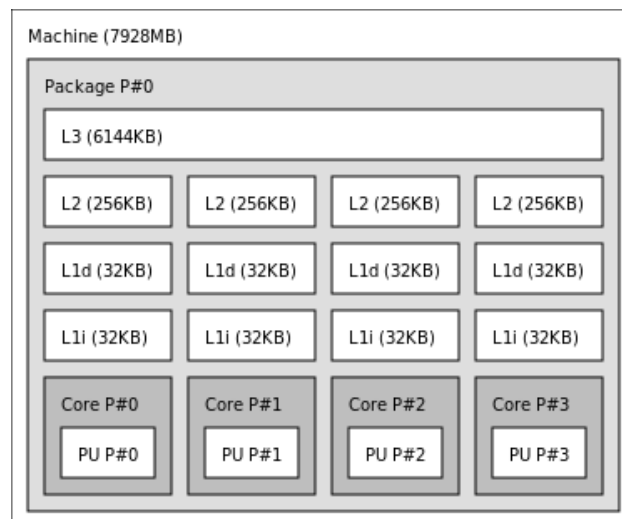   - L3: 12-way set associative

Figure 1: Topology information of a Dell OptiPlex 7020

1.3   Timing Results

In the following tables and plots the recorded execution times are displayed. Note that X axis is plotted on a (base 2) <u>logarithmic</u> scale while Y axis on a <u>linear</u> scale.

Timing results of π calculation (Time unit: seconds)						
Chunk Size	# of threads	1st run	2nd run	3rd run	4th run	Average time
1	1	18.451202	18.443870	18.444319	18.441278	18.44516725
1	4	98.559317	98.393137	99.515415	98.189223	98.664273
1	16	95.482310	95.205719	95.275233	95.197046	95.290077

Table 1: Timing results of π calculation using chunk size = 1 iteration

Figure 2: Timing results of π calculation using chunk size = 1 iteration

Timing results of π calculation (Time unit: seconds)						
Chunk Size	# of threads	1st run	2nd run	3rd run	4th run	Average time
10	1	6.505206	6.510850	6.507051	6.511070	6.50854425
10	4	10.843631	10.728116	10.715714	10.832101	10.7798905
10	16	10.829372	10.820372	10.842818	10.748566	10.810282

Table 2: Timing results of π calculation using chunk size = 10 iterations

Figure 3: Timing results of π calculation using chunk size = 10 iterations

Timing results of π calculation (Time unit: seconds)						
Chunk Size	# of threads	1st run	2nd run	3rd run	4th run	Average time
100	1	6.275921	6.279012	7.893470	6.281098	6.68237525
100	4	2.428611	2.464799	2.463414	2.425332	2.445539
100	16	2.425184	2.459710	2.432488	2.458897	2.44406975

Table 3: Timing results of π calculation using chunk size = 10^2 iterations

Figure 4: Timing results of π calculation using chunk size = 10^2 iterations

| Timing results of π calculation (Time unit: seconds)						
Chunk Size	# of threads	1st run	2nd run	3rd run	4th run	Average time
1000	1	6.248489	6.254362	6.254913	6.251492	6.252314
1000	4	1.733363	1.735331	1.732440	1.734742	1.733969
1000	16	1.731060	1.726669	1.730891	1.732937	1.73038925

Table 4: Timing results of π calculation using chunk size $= 10^3$ iterations

Figure 5: Timing results of π calculation using chunk size $= 10^3$ iterations

Timing results of π calculation (Time unit: seconds)						
Chunk Size	# of threads	1st run	2nd run	3rd run	4th run	Average time
10000	1	6.244337	6.252478	6.250002	6.252064	6.24972025
10000	4	1.664214	1.659239	1.660260	1.659163	1.660719
10000	16	1.660762	1.664066	1.661164	1.658869	1.66121525

Table 5: Timing results of π calculation using chunk size $= 10^4$ iterations

Timing results of π calculation (Time unit: seconds)						
Chunk Size	# of threads	1st run	2nd run	3rd run	4th run	Average time
100000	1	6.237799	6.234975	6.244593	6.235083	6.2381125
100000	4	1.661888	1.658459	1.667569	1.651900	1.659954
100000	16	1.653965	1.652250	1.651300	1.651015	1.6521325

Table 6: Timing results of π calculation using chunk size $= 10^5$ iterations

Figure 6: Timing results of π calculation using chunk size $= 10^4$ iterations

Figure 7: Timing results of π calculation using chunk size $= 10^5$ iterations

1.4   Conclusion

Based on the results presented above and given that the average execution time of the serial program is 6,245896 seconds, we conclude that:

- Program performance is increased until oversubscription appears. Even though we expect overheads to be introduced due to time slicing (e.g. context switching, cache pollution), the execution time becomes approximately constant after the number of threads exceeds the number of the processors available (4). This happens because switching between threads is less resource-intensive than switching between processes.

- Self-scheduling leads to execution times multiple times greater than the one of the serial program and in case of multithreaded calculation, dozens of times greater. The chunk size of one (1) iteration is a fine-grained task something that results in threads constantly racing to acquire the same mutex lock. As the number of threads is increased, race overhead is increased too. Moreover, the granularity of self-scheduling leads to more function calls taking place, something that adds up to the existing overheads.

- A single thread executing tasks with chunk size ≥ 10 requires almost the same time as the serial program.

- Multiple threads and tasks with chunk size $\ll 10^2$ result in higher execution times compared to the serial program. The reasons for these overheads are the same as in self-scheduling (fine-grained parallelism).

- Parallel program efficiency is unfolded for chunk size $\geq 10^2$ but hits a bottleneck for more coarse-grained tasks (chunk size $\geq 10^4$ iterations in this case).

2   Exercise #2

2.1   About

This exercise is about the multiplication of integer NxN arrays using POSIX threads and static scheduling. During static scheduling the parallelizable loops are evenly (when possible) divided into chunks of iterations (tasks) and are dispatched to the threads available to the runtime system for execution. Due to this even distribution of iterations and in contrast to dynamic scheduling, a thread executing on a processor under heavy workload will have a lower throughput capability and will lead to an increase of the total execution time. The purpose of this exercise is to parallelize only the <u>outermost</u> for-loop of the serial calculation, time the matrix multiplication and observe how altering the number of threads will affect execution time.

2.2   Implementation details

If T is the number of threads, and N is the array dimension, then the outermost for-loop of the serial program consists of N iterations, that should be divided into T chunks. When T divides N evenly, the exact chunk size is N/T. In the opposite case, $S = N \bmod T$ threads will be assigned $\lfloor N/T \rfloor + 1$ iterations and $T - S$ threads will be assigned $\lfloor N/T \rfloor$ iterations. We are going to refer to S as the number of special threads because these threads execute one more iteration than the rest. This policy manages to avoid a lopsided distribution of iterations to threads as it may increase workload by only a single iteration [1].

---

[1]This additional iteration may add significant delays in coarse-grained tasks

2.3 Experiment details

During this experiment, thread number takes value in {1, 2, 4, 8, 12, 16} and array size is 1024x1024.

2.3.1 System Specifications

The experiments were conducted again on a Dell OptiPlex 7020.

2.4 Timing Results

In the following table and plot the recorded execution times are displayed. Note that X axis is plotted on a (base 2) logarithmic scale while Y axis on a linear scale.

Timing results of matrix multiplication (Time unit: seconds)				
Array size: 1024x1024				
# of threads	1st run	2nd run	3rd run	4th run	Average time
1	5.057095	4.319662	3.157559	5.297035	4.45783775
2	2.648351	2.545161	1.940197	2.043663	2.294343
4	0.918773	0.800195	1.504583	1.504583	1.07884625
8	0.971947	1.428921	1.224236	0.990774	1.1539695
12	0.965793	0.986273	0.981548	1.130868	1.0161205
16	1.208460	1.241845	1.015431	0.974093	1.10995725

Table 7: Timing results of 2D matrix multiplication

Figure 8: Timing results of 2D matrix multiplication

2.5   Conclusion

Based on the results presented above and given that the average execution time of the serial program is 4,297561 seconds, we conclude that:

- Program performance is approximately doubled as the number of threads is increased, until oversubscription bottleneck is hit and execution time becomes almost constant [2].

- Execution time of the parallel program when a single thread is used to perform the calculation is a bit higher than the time of the serial program. Moreover, we observed that execution time in general varies from time to time. Several reasons such as context switching, thread affinity, cache misses & cache pollution justify the existence of these overheads and consequently this behavior.

3   Exercise #3

3.1   About

This exercise is about implementing a (simple) custom barrier for a group of POSIX threads using only integer data types and POSIX condition variables [3]. The purpose of this exercise is to observe how altering the number of threads will affect execution time of user programs using custom barrier synchronization and how it is compared to the execution time while using the POSIX threads' barrier implementation.

3.2   Experiment details

During this experiment, thread number takes value in {1, 2, 4, 8, 32, 128, 512, 1024}.

3.2.1   System Specifications

The experiments were conducted once again on a Dell OptiPlex 7020.

3.3   Implementation Details

Implementing the barrier synchronization mechanism required definition of struct barrier_s and of the following three functions:

```
int barrier_init(barrier_t *barrier, unsigned int nthr);
int barrier_wait(barrier_t *barrier);
int barrier_destroy(barrier_t *barrier);
```

Listing 1: Struct barrier_s

```
0  typedef struct barrier_s {
1      unsigned int     init_count;
2      unsigned int     arrived;
3      unsigned int     left;
4      pthread_cond_t  *release_threads;
5      pthread_cond_t  *next_bar;
6      pthread_mutex_t *mutex;
7  } barrier_t;
```

---

[2]This is also the conclusion we came to on Exercise #1

[3]The use of global variables is not allowed

3.3.1 Struct `barrier_s`

Struct `barrier_s` contains all the info required to implement the custom barrier and consists of the following members:

- `init_count` : unsigned int
  The number of threads that must call `barrier_wait()` before any of them return to the caller.

- `arrived` : unsigned int
  The number of threads that have arrived to the barrier; Have called `barrier_wait()` and are currently blocked.

- `left` : unsigned int
  The number of threads that have left the barrier; Returned from the `barrier_wait()` call.

- `release_threads` : pthread_cond_t
  The threads arriving at the barrier block on this condition variable until all `init_count` threads arrive.

- `next_bar` : pthread_cond_t
  The threads arriving at the barrier, while threads from a previous phase are currently leaving the barrier, block on this condition variable until all currently blocked threads have left and the barrier can be reused.

3.3.2 `barrier_init()`

The `barrier_init()` function allocates the resources required to use the barrier referenced by <u>barrier</u> and initializes them as needed. The implementation of this function is pretty straightforward so I am not going to further explain how it works. The results are undefined if this function is called when any thread is blocked on the barrier or <u>barrier</u> is not initilized. If `barrier_init()` function fails, the contents of the barrier are undefined. Upon successful completion, this function returns zero unless one of the following errors occurs:

- `EINVAL`: The value specified by <u>count</u> is equal to zero or <u>barrier</u> is `NULL`.

- `ENOMEM`: Insufficient memory exists to initialize the barrier.

- `EBUSY`: If the implementation detects that the <u>barrier</u> argument refers to an already initialized barrier object.

3.3.3 `barrier_destroy()`

The `barrier_destroy()` function destroys the barrier referenced by <u>barrier</u> and releases the resources it currently holds. The implementation of this function is pretty straightforward too so I am not going to further explain how it works either. Use of the <u>barrier</u> after calling `barrier_destroy()` is undefined. The results are undefined if this function is called when any thread is blocked on the barrier or <u>barrier</u> is not initilized. Upon successful completion, this function returns zero unless the following error occurs:

- `EINVAL`: The barrier object referenced by <u>barrier</u> is `NULL`.

3.3.4 `barrier_wait()`

The `barrier_wait()` function synchronizes participating threads at the barrier referenced by <u>barrier</u>. The calling thread blocks until `init_count`[4] threads have called `barrier_wait()` specifying the very same

---

[4]Has the same value as `count`, specified during `barrier_init()` call.

barrier object. When the required number of threads have arrived at the barrier, all threads are unblocked and the <u>barrier</u> is reset for future usage. The results are undefined if this function is called with an uninitialized barrier. Upon successful completion, this function returns `PTHREAD_BARRIER_SERIAL_THREAD` [5] to a single arbitrary thread synchronized to the barrier and zero to the remaining threads synchronized. Function `barrier_wait()` may fail with the following error:

- `EINVAL`: The barrier object referenced by <u>barrier</u> is `NULL`.

The `barrier_wait(barrier_t *barrier)` function works as follows:

1. If <u>barrier</u> argument is `NULL`, `EINVAL` is returned.

2. The barrier's mutex is locked to ensure that the encountering thread is the only thread around.

3. In case there are threads exiting the barrier, the encountering thread blocks on barrier's condition variable `next_bar`. As soon as it unblocks, it has the mutex lock already acquired and proceeds to step #4 as if this step didn't exist.

4. The value of member `arrived` is incremented.

5. In case not all `init_count` threads have called `barrier_wait()`, the encountering thread blocks on barrier's condition variable `release_threads`. As soon as it unblocks, it has the mutex lock acquired and proceeds to step #6.

6. The first thread that will reach this part of the code is the last thread that called `barrier_wait()`. This happens because only that thead will not block on condition variable `release_threads` and will later signal the rest `init_count−1` threads.

7. The value of member `left` is incremented.

8. If the encountering thread is the <u>first</u> leaving the barrier, it unblocks all threads blocked on `release_threads` condition variable and the function's return value is set to `PTHREAD_BARRIER_SERIAL_THREAD`.

9. If the encountering thread is the <u>last</u> one leaving the barrier, it resets `left` and `arrived` members to zero and unblocks any threads blocked on `next_bar` condition variable as the barrier is ready for reuse.

10. The barrier's mutex is unlocked.

11. Zero is returned unless the return value has been set to `PTHREAD_BARRIER_SERIAL_THREAD`.

As you may have observed, the above implementation avoids deadlocks.

3.4   Bugs

Concurrently calling `barrier_init()` or `barrier_destroy()` has undefined results. These two functions should be atomically executed using low level synchronization mechanisms. Perhaps Linux Kernel Futexes[6] should do the job.

---

[5] Constant defined in `pthread.h`
[6] Fast Userspace muTEX-es

3.5 Timing Results

In the following table and plot the recorded execution times are displayed. Note that X axis is plotted on a (base 2) <u>logarithmic</u> scale while Y axis on a (base 10) <u>logarithmic</u> scale.

Timing results of program execution (Time unit: miliseconds)								
	# of threads							
Barrier Implemetation	1	2	4	8	32	128	512	1024
Custom	0.125	0.69325	1.56775	3.952	30.27725	119.7172	431.333	950.200
POSIX threads	0.085	1.09525	1.84725	3.0225	16.98425	76.451	298.4285	643.386

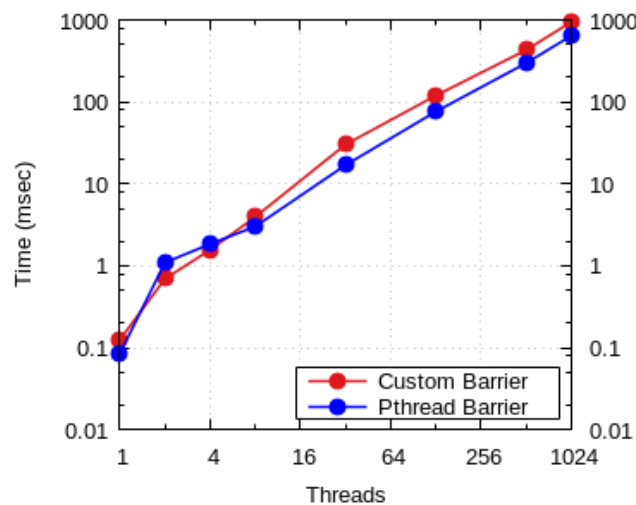Table 8: Timing results of program using barrier synchronization

Figure 9: Timing results of program using barrier synchronization

3.6 Conclusion

Based on the results presented above we conclude that:

- User program execution time has a near exponential growth relative to the number of threads synchronized by the barrier. This is observed for both the custom and the POSIX barrier implementation.

- The custom barrier implementation results in higher execution times. This happens for the following reasons:

  1. The custom barrier implementation is not as optimized as possible because I implemented an algorithm I came up with during this assignment[7].

  2. The custom barrier implementation uses POSIX threads' condition variables and a mutex, while the POSIX threads' barrier uses calls to low level synchronization mechanisms found in glibc library[8].

---

[7]To probe further check Sense-Reversing Barriers
[8]For more information refer to lowlevellock.h